

Project 1: Remote Procedure Calls

Due: September 30th, 2014 23:59:59pm

In this project, you will create a simple Remote Procedure Call service using Apache Thrift.

I have provided a file, `fileservice.thrift`, that defines several basic operations that can be conducted on a remote file server.

The project consists of three parts, which are described below.

1 Compile the interface definition file

I have provided an Apache Thrift interface definition file, `fileservice.thrift`. It can be downloaded from Blackboard. You can use Thrift to compile interface definition file to either Java, C++, or Python (Thrift also supports other languages, but these are not available on CS computers).

I have compiled and installed Thrift under my home directory on CS Department computers. These computers are located in G7 and Q22, or you can also remotely access them by SSH'ing into `remote.cs.binghamton.edu`. Source code as well as compiled files are located under `/home/yaoliu/src_code/thrift`. To use Thrift installed under my home directory, you need to set the environment variable `PATH`:

```
$> export PATH=$PATH:/home/yaoliu/src_code/local/bin
```

It is strongly encouraged that you install Thrift under your home directory or on your personal computer. You can download the Thrift source code at <http://thrift.apache.org/download>. To compile the source code and install, follow the instructions listed in <http://thrift.apache.org/docs/BuildingFromSource>.

The following commands can be used to install Thrift under your home directory on CS computers (but may not work on other system configurations):

```
$> git clone https://git-wip-us.apache.org/repos/asf/thrift.git thrift
$> cd thrift
```

```
$> ./bootstrap.sh
$> ./configure --prefix=$HOME/local --with-haskell=no
```

```
$> make
$> make install
```

```
$> export PATH=$PATH:$HOME/local/bin
```

Thrift can be used to create service stubs in a language of your choice from a `.thrift` interface definition file. For information on how to do so, refer to <http://thrift.apache.org/tutorial>. As an example, the Thrift command to compile this project's IDL file to Python is as follows:

```
$> thrift -gen py fileservice.thrift
```

In <http://thrift.apache.org/tutorial>, you can also find example clients and servers written in many of the languages that Thrift supports. This example code can be found in the Thrift source tree under the `tutorial` directory. I recommend that you read the tutorial code in the language of your choice before beginning work on the remainder of the project.

Information about the basic concepts and features of Thrift are available at <http://thrift.apache.org/docs/concepts> and <http://thrift.apache.org/docs/features>. Although I have provided the IDL file for you, you should also check <http://thrift.apache.org/docs/idl> for details about the Thrift interface definition language and data types and services supported by Thrift.

2 Extend the server-side method stubs generated by Thrift

When Thrift processes the interface definition file, `fileservice.thrift`, it generates method stubs corresponding to the defined services. You must implement methods corresponding to the following server-side stubs:

writeFile given a name, owner, and contents, the corresponding file should be written to the server. Meta-information, such as the owner, creation-time, update-time, version, content length, and content hash (use the MD5 hash) should also be stored at the server side. If the filename does not exist on the server, a new file should be created with its version attribute set to 0. Otherwise, the file contents should be overwritten and the version number should be incremented.

readFile if a file with a given name and owner exists on the server, both the contents and meta-information should be returned. Otherwise, an exception should be thrown.

deleteFile if a file with a given name and owner exists on the server, it should be deleted, but the metadata should be retained and the deleted timestamp modified to record the time of deletion. If an error occurs during file deletion, e.g., deleting a non-existent file, an exception should be thrown.

listOwnerFiles given an user name, all files owned by the user should be listed. Files that have been deleted should not be listed. If the user owns no files, then an empty list should be returned.

Every time the server is started any filesystem or other persistent storage that it used should be initialized to be empty. If your server stores files in the file system, it should use only the current working directory from where it was run. There is no need to implement directory structure.

The server executable should take a single command-line argument specifying the port where the Thrift service will listen for remote clients (The default Thrift server port is 9090). It should also use Thrift's `TBinaryProtocol` for marshalling and unmarshalling data structures.

3 Write a client program

You must also write a client program to call the remote methods that you created in Part 2. The client program should take five command line arguments as input. The first two arguments should be the host and port of the remote Thrift server. The final three arguments (occurring in no particular order) are as follows:

--operation [operation], where operation can be "read", "write", "delete", and "list", and should invoke the respective remote procedure call

--filename [filename], should be used as input to specify either a local file to be written to the remote system, or a remote filename to be read or deleted

--user [user], should specify the owner of the file for the current operation

After the client program receives the response from the remote procedure call, it should use a separate Thrift transport and protocol to write the returned structure to the standard output. (Hint: Under Python, use TFileObjectTransport and TJSONProtocol. Under Java, use TIOStreamTransport and TJSONProtocol.) **This output step is important!** If it is not correctly implemented, then the automated test scripts used to grade your work will fail, causing you to receive a low grade on the project. The client should also write any exceptions thrown by the remote method in JSON format to the standard output.

An example of how the client program should operate is provided below:

Input:

```
$> ./client localhost 9090 --operation read --filename example.txt --user guest
```

The output is in the form of a Thrift-generated JSON¹ structure (do not try to manually generate this yourself):

```
{"1":{"rec":{"1":{"str":"example.txt"},"2":{"i64":1379202165425},
"3":{"i64":1379202165425},"5":{"i32":0},"6":{"str":"guest"},"7":{"i32":10},
"8":{"str":"7bdbc2db770bf3dc4f78ab4fbf88831a"}}},"2":{"str":"example contents\n"}}
```

Note: When accessing `remote.cs.binghamton.edu`, you are actually redirected to one of the 16 REMOTE machines (`{remote00, remote01, ..., remote09, remote10, ..., remote15}.cs.binghamton.edu`) using DNS redirection. So to test your implementation, you need to make sure your client is contacting the server using the correct host name or public IP address.

4 How to submit

To submit the project, you should first create a directory whose name is "your BU email ID"-project1. For example, if your email ID is `jdoue@binghamton.edu`, you should create a directory called `jdoue-project1`. You should put the following files into this directory:

1. The `fileservice.thrift` file I provided.
2. Your source code which includes at least two files implementing the client and the server, respectively.
3. A Makefile to compile your source code into two executables, which should be named `client` and `server`. (It is okay if these executables are bash scripts that call the Java interpreter, as long as the command line arguments follow the format described in Parts 2 & 3.)
4. A Readme file describing the programming language(s) you are using, your implementation details, and sample input/output.

Compress the directory (e.g., `tar czvf jdoue-project1.tar.gz jdoue-project1`) and submit the tarball (in `tar.gz` or `tgz` format) to Blackboard. Again, if your email ID is `jdoue@binghamton.edu`, you should name your submission: `jdoue-project1.tar.gz` or `jdoue-project1.tgz`. Failure to use the right naming scheme will result in a 5% point deduction.

Your project will be graded on the CS Department computers `remote.cs.binghamton.edu`. If you use external libraries in your code, you should also include them in your directory and correctly set the environment variables using absolute path in the Makefile. If your code does not compile on or cannot correctly run on the CS computers, you will receive no points.

¹<http://en.wikipedia.org/wiki/JSON>